

Student: Duncan Ferguson

Student Id: 871641260

Class: Comp 4431-1

Assignment: Assignment 2

Date: 11/12/2021

Project was Completed by Individual, group is just listed below for reference to past code that might be reused from previous Exercises

Group: Name: Broken Toe

Group Members: Emma Bright, Mike Santoro

Assignment Description:

In this assignment you are to use Decision trees and Naive Bayes for classification. In addition, you are to explore your models to determine which feature vector attributes you can remove without "significantly hurting" classification importance. There is no single "right" answer for this question. Also, "significantly hurting" is a judgement call, please explain why you mean by that and provided numeric values. You should start with the full dataset and look at the decision tree `features_importances_` and naive bayes `permutation_importance` to guide which attributes to remove. You will then want to iteratively remove attributes, possibly one at a time or several at a time. Again, there is no single "right way" to do this.

The goal is to predict whether it will rain tomorrow given the days weather statistics. The last column, "Rain Tomorrow" is the classification value for I have cleaned the data set for you. I removed a few attributes that were almost all zeros, and replaced missing values with the mean value for that column. Also, strings have been transformed into numeric values so the data set will work well for both decision trees and naive bayes.

What To Turn In:

1). A report that explains which attributes you have eliminated, why those, and what effect this has on the quality of predictions compared to using the full dataset. Feel free to include a table of `feature_importances` and `permutation_importances` that show how any why you reduced the set of attributes.

2). Your code for creating the decision tree and naive bayes models and getting the feature/permutation importances. Your code does not need to cover your whole experimental space, but should provide enough information to show how you made at least one decision about which attribute (column) to eliminate

```
556 import numpy as np
import pandas as pd
```

```

from sklearn import tree
from sklearn.naive_bayes import *
from sklearn.model_selection import train_test_split
from sklearn.metrics import *
from sklearn.inspection import permutation_importance
import matplotlib.pyplot as plt
import seaborn as sns

```

To start of this project I have decided to look at the correlations first to get a bit more of a grasp on how the data is related.

557 # Importing the Data

```


np.random.seed(0)
df = pd.read_csv("assignment2_cleanInfile.csv")
df.head()
df_did_rain = df[df["RainTomorrow"] != 0]
df_did_not_rain = df[df["RainTomorrow"] == 0]
print(df_did_rain["RainTomorrow"].unique())
print(df_did_not_rain["RainTomorrow"].unique())

[1]
[0]

```

558 sns.heatmap(df.corr())

558 <AxesSubplot:>



This Heat Map Correlation gives us some middle of the ground results. We can see that MinTemp and Temp9am are related, as well as with the max temp.

559 corr_matrix = df.corr().abs()
corr_matrix

559

	Location	MinTemp	MaxTemp	Rainfall	WindGustDir	WindGustSpeed	Wir
--	----------	---------	---------	----------	-------------	---------------	-----

	Location	MinTemp	MaxTemp	Rainfall	WindGustDir	WindGustSpeed	Win
Location	1.000000	0.006194	0.020490	0.003457	0.075197	0.069269	0.00
MinTemp	0.006194	1.000000	0.733919	0.103314	0.149557	0.173331	0.00
MaxTemp	0.020490	0.733919	1.000000	0.074202	0.176500	0.066329	0.18
Rainfall	0.003457	0.103314	0.074202	1.000000	0.033045	0.127250	0.08
WindGustDir	0.075197	0.149557	0.176500	0.033045	1.000000	0.129199	0.34
WindGustSpeed	0.069269	0.173331	0.066329	0.127250	0.129199	1.000000	0.10
WindDir9am	0.001938	0.002464	0.181369	0.083133	0.346856	0.105557	1.00
WindDir3pm	0.030553	0.145183	0.158061	0.042696	0.565715	0.142515	0.30
WindSpeed9am	0.077038	0.174946	0.014586	0.085977	0.063066	0.577864	0.18
WindSpeed3pm	0.064180	0.174187	0.050381	0.056762	0.111046	0.658377	0.14
Humidity9am	0.002065	0.232372	0.499777	0.221392	0.031133	0.209272	0.04
Humidity3pm	0.011066	0.005913	0.499725	0.249609	0.009683	0.025822	0.13
Pressure9am	0.036492	0.424357	0.309093	0.159676	0.120842	0.426600	0.06
Pressure3pm	0.046345	0.434034	0.397422	0.120366	0.042960	0.384642	0.01
Cloud9am	0.010908	0.061460	0.226776	0.171613	0.054864	0.052281	0.07
Cloud3pm	0.015860	0.015902	0.214469	0.145848	0.054884	0.080202	0.04
Temp9am	0.015596	0.897999	0.880087	0.011384	0.169948	0.146756	0.07
Temp3pm	0.022718	0.699828	0.969735	0.077553	0.188980	0.032274	0.19
RainToday	0.004146	0.055856	0.227421	0.500997	0.104836	0.148977	0.16
RainTomorrow	0.003579	0.083717	0.159087	0.236874	0.032470	0.225264	0.03

```

560 def train_data(df, drops=None):
    """This Function splits the data to train and test the data. It also an element for being able to
    drop certain rows."""
    # Y is the classification
    Y = df['RainTomorrow'].tolist()
    X = df.copy()
    X = X.drop(columns=["RainTomorrow"])
    if drops != None:
        X = X.drop(columns=drops)
    x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.25, random_state=42)
    return x_train, x_test, y_train, y_test, X

561 def decisionTree_Model(x_train, x_test, y_train, y_test, X, yes_print=False):
    """This Function goes through building a decision tree and creating a confusion matrix
    computing accuracy, and listing the decision tree features and their importance"""
    # Building out a decision tree

```

```

dtree = tree.DecisionTreeClassifier(criterion="gini")
dtree = dtree.fit(x_train, y_train)
y_predicted = dtree.predict(x_test)
accuracy = accuracy_score(y_test, y_predicted)
important = dtree.feature_importances_
df_importance_list = []
for i, v in enumerate(important):
    df_importance_list.append([X.columns[i], v])
df_importance = pd.DataFrame(df_importance_list, columns=["FName", "Score"])
df_importance.sort_values(by=['Score'], ascending=False, inplace=True)
decisionTree_Model_Features = df_importance["FName"].tolist()
if yes_print:
    print("DecisionTree Accuracy: ", accuracy)
    print("DecisionTree Confusion Matrix:")
    print(confusion_matrix(y_test, y_predicted))
    print("decision tree dtree feature importance:")
    print(df_importance)
return accuracy, decisionTree_Model_Features

```

```

562 def GaussianNB_Model(x_train, x_test, y_train, y_test, X, yes_print=False):
    """This Function goes through the GaussianNB model and calculates the accuracy, computes the confusion
    and then lists the Gaussian Model Features and their significance"""
    model2 = GaussianNB()
    model2.fit(x_train, y_train)
    gaussianNB_predicted = model2.predict(x_test)
    accuracy = accuracy_score(y_test, gaussianNB_predicted)
    imps = permutation_importance(model2, x_test, y_test)
    df_Gaussian_Feature_Importance_List = []
    Gaussian_Feature_Importance_List = imps.importances_mean.tolist()
    for row in enumerate(Gaussian_Feature_Importance_List):
        df_Gaussian_Feature_Importance_List.append([X.columns[row[0]], row[1]])
    df_Gaussian_Feature_Importance = pd.DataFrame(df_Gaussian_Feature_Importance_List, columns=["Feature",
    df_Gaussian_Feature_Importance.sort_values(by=["Significance"], ascending=False, inplace=True)
    GaussianND_Features = df_Gaussian_Feature_Importance["Feature"].tolist()
    if yes_print:
        print('\nconfusion_matrix from Gaussian naive bayes:')
        print(confusion_matrix(y_test, gaussianNB_predicted))
        print('accuracy = ' + str(accuracy))
        print("gaussianNB feature importance:")
        print(df_Gaussian_Feature_Importance)
    return accuracy, GaussianND_Features

```

```

563 df = pd.read_csv("assignment2_cleanInfile.csv")

x_train, x_test, y_train, y_test, X = train_data(df)
dt_a, dt_f = decisionTree_Model(x_train, x_test, y_train, y_test, X, True)

```

```

# train_data(df)
gm_a, gm_f = GaussianNB_Model(x_train, x_test, y_train, y_test, X, True)

```

```

DecisionTree Accuracy: 0.7843258600804524
DecisionTree Confusion Matrix:
[[23618  4005]
 [ 3662  4264]]
decision tree dtree feature importance:

```

	FName	Score
11	Humidity3pm	0.268860
13	Pressure3pm	0.072560
3	Rainfall	0.063940
5	WindGustSpeed	0.060039
1	MinTemp	0.050088

```

10 Humidity9am 0.048687
12 Pressure9am 0.047536
16 Temp9am 0.045614
17 Temp3pm 0.045001
0 Location 0.040346
2 MaxTemp 0.039299
9 WindSpeed3pm 0.036592
8 WindSpeed9am 0.034055
7 WindDir3pm 0.031931
4 WindGustDir 0.031225
15 Cloud3pm 0.030189
6 WindDir9am 0.029780
14 Cloud9am 0.021221
18 RainToday 0.003038

```

confusion_matrix from Gaussian naive bayes:

```

[[24213  3410]
 [ 3488  4438]]

```

accuracy = 0.8059579735013643

gaussianNB feature importance:

	Feature	Significance
11	Humidity3pm	0.013919
3	Rainfall	0.013587
13	Pressure3pm	0.005159
5	WindGustSpeed	0.004169
12	Pressure9am	0.003713
15	Cloud3pm	0.002729
1	MinTemp	0.001474
14	Cloud9am	0.000090
10	Humidity9am	0.000090
8	WindSpeed9am	0.000011
0	Location	-0.000096
9	WindSpeed3pm	-0.000191
7	WindDir3pm	-0.000326
4	WindGustDir	-0.000433
16	Temp9am	-0.000624
6	WindDir9am	-0.000805
17	Temp3pm	-0.001210
2	MaxTemp	-0.001536
18	RainToday	-0.002127

The Code above displays confusion matrix and accuracy of both the Decision Tree Model and the Gaussian Naive Bayes models with all of the features included. In this instance the Gaussian model was more accurate than the decision tree model. Both models show that rain today has the lowest score/significance in the models. One would assume that to make these models more accurate we would remove "RainToday" to improve the models.

I have set up the train_data function to drop features for retraining the data to be run back through the models.

```

564 # Dropping Rain Today
x_train, x_test, y_train, y_test, X = train_data(df, "RainToday")
# Decision Tree Model
dt_a, dt_f = decisionTree_Model(x_train, x_test, y_train, y_test, X, True)

DecisionTree Accuracy: 0.7815128414301387
DecisionTree Confusion Matrix:
[[23533  4090]
 [ 3677  4249]]
decision tree dtree feature importance:

```

	FName	Score
11	Humidity3pm	0.267915
13	Pressure3pm	0.073500
3	Rainfall	0.065227
5	WindGustSpeed	0.060248
1	MinTemp	0.053279
10	Humidity9am	0.049318
12	Pressure9am	0.047377
16	Temp9am	0.046591
17	Temp3pm	0.043340
0	Location	0.040075
2	MaxTemp	0.039315
9	WindSpeed3pm	0.037006
8	WindSpeed9am	0.034378
7	WindDir3pm	0.031334
4	WindGustDir	0.030908
15	Cloud3pm	0.030401
6	WindDir9am	0.029767
14	Cloud9am	0.020021

```
565 # Gaussian Naive Bayes Model
gm_a, gm_f = GaussianNB_Model(x_train, x_test, y_train, y_test, X, True)
```

confusion_matrix from Gaussian naive bayes:

```
[[24720 2903]
 [ 3617 4309]]
```

accuracy = 0.8165911839995499

gaussianNB feature importance:

	Feature	Significance
11	Humidity3pm	0.023117
3	Rainfall	0.017109
13	Pressure3pm	0.005840
5	WindGustSpeed	0.005592
12	Pressure9am	0.004923
15	Cloud3pm	0.003567
10	Humidity9am	0.003483
1	MinTemp	0.002627
14	Cloud9am	0.000647
0	Location	-0.000304
4	WindGustDir	-0.000636
16	Temp9am	-0.000647
7	WindDir3pm	-0.000720
6	WindDir9am	-0.000731
8	WindSpeed9am	-0.000793
9	WindSpeed3pm	-0.000866
17	Temp3pm	-0.001243
2	MaxTemp	-0.001542

Now looking at both the models with Raintoday removed we notice something interesting. The accuracy for the Gaussian model went up. While the Accuracy for the Decision Tree went slightly down.

The Next Step is to work with the Gaussian model removing features that have a negative significance. I have set up the code so that we can look at the accuracy, followed by the features that are included

```
566 x_train, x_test, y_train, y_test, X = train_data(df, ["RainToday", "MaxTemp"])
gm_a, gm_f = GaussianNB_Model(x_train, x_test, y_train, y_test, X, True)
```

$$\begin{bmatrix} [24901 & 2722] \\ [3705 & 4221] \end{bmatrix}$$

```
gaussianNB feature importance:
```

	Feature	Significance
10	Humidity3pm	0.024124
2	Rainfall	0.018369
14	Cloud3pm	0.005272
9	Humidity9am	0.004912
4	WindGustSpeed	0.004518
12	Pressure3pm	0.003972
11	Pressure9am	0.002729
16	Temp3pm	0.001660
13	Cloud9am	0.001384
1	MinTemp	0.001170
15	Temp9am	0.000281
0	Location	0.000236
6	WindDir3pm	-0.000006
3	WindGustDir	-0.000051
5	WindDir9am	-0.000411
7	WindSpeed9am	-0.000850
8	WindSpeed3pm	-0.001176

The Gaussian Model continues to improve as we remove the features with the lowest significance until we start to remove features that have a positive significance. The Two following chunks of code show the reversal of accuracy improving.

```
567 x_train, x_test, y_train, y_test, X = train_data(df, ["RainToday",
    "MaxTemp",
    "WindSpeed3pm",
    "WindGustDir",
    "WindDir3pm",
    "WindSpeed9am",
    "WindDir9am",
    "Temp9am",
    "Location"])
gm_a, gm_f = GaussianNB Model(x_train, x_test, y_train, y_test, X, True)
```

[[25081, 2542]]

```
[ 3683  4243]]
```

```
accuracy = 0.8248895890179752
```

gaussianNB feature importance:

	Feature	Significance
4	Humidity3pm	0.026285
1	Rainfall	0.019106
2	WindGustSpeed	0.008945
8	Cloud3pm	0.005232
6	Pressure3pm	0.005137
5	Pressure9am	0.003994
3	Humidity9am	0.003668
9	Temp3pm	0.001890
7	Cloud9am	0.001885
0	MinTemp	0.001710

[illegible]

```

        "WindGustDir",
        "WindDir3pm",
        "WindSpeed9am",
        "WindDir9am",
        "Temp9am",
        "Location",
        "MinTemp"])
gm_a, gm_f = GaussianNB_Model(x_train, x_test, y_train, y_test, X, True)

```

confusion_matrix from Gaussian naive bayes:

```

[[25075  2548]
 [ 3722  4204]]

```

accuracy = 0.823623730625334

gaussianNB feature importance:

	Feature	Significance
3	Humidity3pm	0.024563
0	Rainfall	0.018673
1	WindGustSpeed	0.008996
5	Pressure3pm	0.005182
7	Cloud3pm	0.004670
4	Pressure9am	0.004563
2	Humidity9am	0.002678
8	Temp3pm	0.000675
6	Cloud9am	0.000624

Accuracy stops improving after we remove every feature that has a negative significance value. The two models above show the point at which the model no longer improves when removing variables.

These results show model manipulation using significance. Leaving us with the best model for the Gaussian Naive Bayes Model using backward selection as the model that includes:

Model: Gaussian Naive Bayes

Accuracy: 0.8248895890179752

Features: [Humidity3pm, RainFall, WindGustSpeed, Cloud3pm, Humidity9am, Pressure3pm, Pressure9am, Temp3pm, Cloud9am, Mintemp]

Now to take a look at manipulating the decision tree model. If we note from above that when we removed the lowest score feature, "RainToday", the model actually lost accuracy. Because of this I have decided to code in a backward selection by removing one variable at a time until the accuracy improves.

This process is then repeated with the newly improved model. Until Accuracy no longer improves. To save space in the print out. I have included the first step, followed all the way to the last two steps.

Below is the function that is used for removing a feature at a time. The selection input represents the variables that will be removed from the model.

```

569 def backward_selection(accuracy, features, selection, df, model):
    """This is a backward selection model that removes one feature at a time to see if we can improve the

```



```

accuracy of the model. Once any improvement is reached the function returns which feature that was dr
that improves the model. The features that were included and the accuracy of the new model""
features_loop = [word for word in features if word not in selection]
accuracy_start = accuracy
for feature in features_loop:
    x_train, x_test, y_train, y_test, X = train_data(df, [feature] + selection) # Dropping selection
    if model == 1: # 1 For decisionTree
        new_accuracy, new_features = decisionTree_Model(x_train, x_test, y_train, y_test, X, False)
    else:
        new_accuracy, new_features = GaussianNB_Model(x_train, x_test, y_train, y_test, X, False)
    if new_accuracy >= accuracy_start:
        print("Accuracy:", new_accuracy, "\nFeatures:", new_features, "\nDropped Feature:", feature)
        return
return "No Improvement"

```

```

570 # Starting the backward selection model with no drop features.
    features = df.columns.tolist()
    features.pop(-1)
    accuracy = 0.7815128414301387
    selection = []

    backward_selection(accuracy, features, selection, df, 1)

    Accuracy: 0.7832569129933331
    Features: ['Humidity3pm', 'Pressure3pm', 'Rainfall', 'WindGustSpeed', 'MinTemp', 'Humidity9am', 'Pressure
    Dropped Feature: Location

571 selection = ["Location"]
    accuracy = 0.7832569129933331
    backward_selection(accuracy, features, selection, df, 1)

    Accuracy: 0.7865481448142001
    Features: ['Humidity3pm', 'Pressure3pm', 'WindGustSpeed', 'Rainfall', 'MinTemp', 'Pressure9am', 'Humidity
    Dropped Feature: Cloud3pm

589 selection = ["Location", "Cloud3pm"]
    accuracy = 0.7865481448142001
    backward_selection(accuracy, features, selection, df, 1)

589 'No Improvement'

```

This process shows us that removing the [Location, Cloud3pm] leaves us with the best accuracy on the model.

Process: Backward Selection

Model: Decision Tree

Accuracy: 0.7865481448142001

Features:['Humidity3pm', 'Pressure3pm', 'WindGustSpeed', 'Rainfall', 'MinTemp', 'Pressure9am', 'Humidity9am', 'Temp9am', 'Temp3pm', 'MaxTemp', 'WindSpeed3pm', 'WindSpeed9am', 'WindDir9am', 'WindGustDir', 'WindDir3pm', 'Cloud9am', 'RainToday']

Using the same backward selection model methodology the next steps test the Gaussian Model to see if we can improve the model that we were left with by removing all the features with

negative significance.

```
590 selection = ["RainToday", "MaxTemp", "WindSpeed3pm", "WindGustDir", "WindDir3pm", "WindSpeed9am", "WindDir",  
              "Temp9am", "Location"]  
accuracy = 0.8248895890179752  
backward_selection(accuracy, features, selection, df, 0)
```

```
Accuracy: 0.8256491040535598  
Features: ['Humidity3pm', 'Rainfall', 'WindGustSpeed', 'Cloud3pm', 'Temp3pm', 'Cloud9am', 'Pressure3pm',  
Dropped Feature: Humidity9am
```

```
591 Accuracy= 0.8256491040535598  
selection = ['Humidity3pm', 'Rainfall', 'WindGustSpeed', 'Cloud3pm', 'Temp3pm', 'Cloud9am', 'Pressure3pm',  
            'Pressure9am', 'MinTemp', "Humidity9am"]  
backward_selection(accuracy, features, selection, df, 0)
```

```
591 'No Improvement'
```

Looking at the results above. We can slightly improve the Gaussian model that we were given. By dropping Humidity9am from the model we were able to improve the Accuracy to 0.8256491040535598

Process: Remove negative significance features, then backward selection model

Model: Gaussian Naive Bayes

Accuracy: 0.8256491040535598

Features Included: ['Humidity3pm', 'Rainfall', 'WindGustSpeed', 'Cloud3pm', 'Temp3pm', 'Cloud9am', 'Pressure3pm', 'Pressure9am', 'MinTemp', "Humidity9am"]

The next steps are the Gaussian Naive Bayes model using only Backward Model Selection. As noted, I will remove the steps in the middle to save space.

```
592 selection = []  
accuracy = 0  
backward_selection(accuracy, features, selection, df, 0)
```

```
Accuracy: 0.8059861036878675  
Features: ['Humidity3pm', 'Rainfall', 'Pressure3pm', 'WindGustSpeed', 'Pressure9am', 'Cloud3pm', 'MinTemp',  
Dropped Feature: Location
```

```
587 selection = ["Location", "MaxTemp", "WindGustDir", "WindDir9am", "WindSpeed9am", "MinTemp", "WindDir3pm",  
              "Humidity9am", "Temp9am", "RainToday", "Pressure9am"]  
accuracy = 0.8249177192044783  
backward_selection(accuracy, features, selection, df, 0)
```

```
Accuracy: 0.8276744774817857  
Features: ['Humidity3pm', 'Rainfall', 'Cloud3pm', 'WindGustSpeed', 'Pressure3pm', 'Temp3pm']  
Dropped Feature: Cloud9am
```

```
588 selection = ["Location", "MaxTemp", "WindGustDir", "WindDir9am", "WindSpeed9am", "MinTemp", "WindDir3pm",  
              "Humidity9am", "Temp9am", "RainToday", "Pressure9am", "Cloud9am"]  
accuracy = 0.8276744774817857  
backward_selection(accuracy, features, selection, df, 0)
```

After that Long selection Process we were able to come out with a accuracy of

Process: Backward Selection

Model: Gaussian Naive Bayes

Accuracy: 0.8276744774817857

Included Features: ['Humidity3pm', 'Rainfall', 'WindGustSpeed', 'Cloud3pm', 'Pressure3pm', 'Temp3pm']

This Produces a slight improvement from our last most accurate model. Listed below to show the difference.

Process: Remove negative significance features, then backward selection model

Model: Gaussian Naive Bayes

Accuracy: 0.8256491040535598

Features Included: ['Humidity3pm', 'Rainfall', 'WindGustSpeed', 'Cloud3pm', 'Temp3pm', 'Cloud9am', 'Pressure3pm', 'Pressure9am', 'MinTemp', 'Humidity9am']

The results from the backward model selection with the decision tree models left some room for improvement. It is because of this reason that the next steps will test forward model selection.

Below is the function that is set up for forward model selection.

```

616 def forward_selection(accuracy, features, selection, df, model):
    """This is a backward selection model that removes one feature at a time to see if we can improve the
    accuracy of the model. Once any improvement is reached the function returns which feature that was dr
    that improves the model. The features that were included and the accuracy of the new model"""
    features_loop = [word for word in features if word not in selection]
    best_accuracy = accuracy
    for feature in features_loop:
        drop_features = [word for word in features if word not in [feature] + selection]
        x_train, x_test, y_train, y_test, X = train_data(df, drop_features) # Dropping selection and fea
        if model == 1: # 1 For decisionTree
            new_accuracy, new_features = decisionTree_Model(x_train, x_test, y_train, y_test, X, False)
        else:
            new_accuracy, new_features = GaussianNB_Model(x_train, x_test, y_train, y_test, X, False)
        if new_accuracy >= best_accuracy:
            best_accuracy = new_accuracy
            add_feature = feature
    if best_accuracy > accuracy:
        print("Accuracy:", best_accuracy, "\nAdd Feature:", add_feature, "\nModel Features", [add_fea
    else:
        return "No Improvement"

617 features = df.columns.tolist()
    features.pop(-1)
    accuracy = 0
    selection = []

```

```
forward_selection(accuracy, features, selection, df, 1)
```

```
Accuracy: 0.8219921798081521
```

```
Add Feature: Humidity3pm
```

```
Model Features ['Humidity3pm']
```

```
618 selection = ['Humidity3pm']  
accuracy = 0.8219921798081521  
forward_selection(accuracy, features, selection, df, 1)
```

```
Accuracy: 0.8268868322596978
```

```
Add Feature: Location
```

```
Model Features ['Location', 'Humidity3pm']
```

```
619 selection = ['Location', 'Humidity3pm']  
accuracy = 0.8268868322596978  
forward_selection(accuracy, features, selection, df, 1)
```

```
619 'No Improvement'
```

The Forward Selection with the decision tree has vastly improved our model. We only have two features that are being used, but the accuracy is much better. There are a lot of reasons for this. This appears to show that location and Humidity are the best predictors for rain tomorrow.

Process: Forward Model Selection

Model: Decision tree

Accuracy: 0.8268868322596978

Features: ['Location', 'Humidity3pm']

Next Process is to test the Gaussian Model using forward selection

```
621 accuracy = 0  
selection = []  
forward_selection(accuracy, features, selection, df, 0)
```

```
Accuracy: 0.8205012799234859
```

```
Add Feature: Humidity3pm
```

```
Model Features ['Humidity3pm']
```

```
624 accuracy = 0.8291372471799489  
selection = ['Pressure3pm', 'WindGustSpeed', 'Humidity3pm']  
forward_selection(accuracy, features, selection, df, 0)
```

```
Accuracy: 0.8298967622155334
```

```
Add Feature: Cloud3pm
```

```
Model Features ['Cloud3pm', 'Pressure3pm', 'WindGustSpeed', 'Humidity3pm']
```

```
625 accuracy = 0.8298967622155334  
selection = ['Cloud3pm', 'Pressure3pm', 'WindGustSpeed', 'Humidity3pm']  
forward_selection(accuracy, features, selection, df, 0)
```

```
625 'No Improvement'
```

Process: Forward Model Selection

Model: Gaussian Naive Bayes

Accuracy = 0.8298967622155334

Features: ['Cloud3pm', 'Pressure3pm', 'WindGustSpeed', 'Humidity3pm']

Final Conclusion

The best results that we found for both

Process: Full Model Gaussian Naive Bayes Model: Gaussian Naive Bayes Accuracy:

0.8059579735013643 Features: All

Improved to:

Process: Forward Model Selection

Model: Gaussian Naive Bayes

Accuracy = 0.8298967622155334

Features: ['Cloud3pm', 'Pressure3pm', 'WindGustSpeed', 'Humidity3pm']

Process: Full Model Decision Tree Model: Decision Tree Accuracy: 0.7843258600804524 Features: All

Improved to:

Process: Forward Model Selection

Model: Decision tree

Accuracy: 0.8268868322596978

Features: ['Location', 'Humidity3pm']

First we looked at the Significance and the score values for the Gaussian and the Decision Tree. The Significance Values that were negative for the Gaussian model did have significant impact on decreasing the model accuracy. By removing them we were able to improve the accuracy score. When looking at the decision tree feature scores removing the lowest one did not necessarily improve the model.

To expand on these models we used both backward and forward model selection. Backward model selection process removed one features at a time until the accuracy was improved. It then proceeded by removing that feature and then looping all the other features until a greater accuracy was achieved.

The forward model selection added one feature at a time and took the feature that improved the accuracy the most, until adding features no longer improved the model.

The forward model selection produced the greatest improvements in accuracy on both the Decision tree and the Gaussian full models.

